

Does Location influence startup success

Reading the CSV

```
df <- read.csv("C:/Users/chaki/Downloads/cleaned_investments_VC.csv")
```

```
head(df)
```

```
##           name      market funding_total_usd  status
## 1      #waywire      News      1750000  acquired
## 2 'Rock' Your Paper Publishing      40000  operating
## 3 (In)Touch Network Electronics 1500000  operating
## 4 -R- Ranch and Mine      Tourism      60000  operating
## 5 004 Technologies      Software      0  operating
## 6 1,2,3 Listo E-Commerce      40000  operating
##           region      city funding_rounds founded_at founded_month
## 1      New York City  New York      1 2012-06-01      2012-06
## 2      Tallinn      Tallinn      1 2012-10-26      2012-10
## 3      London      London      1 2011-04-01      2011-04
## 4      Dallas Fort Worth      2 2014-01-01      2014-01
## 5 Springfield, Illinois Champaign      1 2010-01-01      2010-01
## 6      Santiago Las Condes      1 2012-01-01      2012-01
## founded_quarter founded_year first_funding_at last_funding_at  seed venture
## 1      2012-Q2      2012      2012-06-30      2012-06-30 1750000      0
## 2      2012-Q4      2012      2012-08-09      2012-08-09  40000      0
## 3      2011-Q2      2011      2011-04-01      2011-04-01 1500000      0
## 4      2014-Q1      2014      2014-08-17      2014-09-26      0      0
## 5      2010-Q1      2010      2014-07-24      2014-07-24      0      0
## 6      2012-Q1      2012      2013-02-18      2013-02-18  40000      0
## equity_crowdfunding undisclosed convertible_note debt_financing angel grant
## 1      0      0      0      0      0      0      0
## 2      0      0      0      0      0      0      0
## 3      0      0      0      0      0      0      0
## 4      60000      0      0      0      0      0      0
## 5      0      0      0      0      0      0      0
## 6      0      0      0      0      0      0      0
## private_equity post_ipo_equity post_ipo_debt secondary_market
## 1      0      0      0      0
## 2      0      0      0      0
## 3      0      0      0      0
## 4      0      0      0      0
## 5      0      0      0      0
## 6      0      0      0      0
## product_crowdfunding round_A round_B round_C round_D round_E round_F round_G
## 1      0      0      0      0      0      0      0      0
## 2      0      0      0      0      0      0      0      0
```

```
## 3      0      0      0      0      0      0      0      0
## 4      0      0      0      0      0      0      0      0
## 5      0      0      0      0      0      0      0      0
## 6      0      0      0      0      0      0      0      0
##   round_H    country age success_metric_updated
## 1      0 United States  2      Successful
## 2      0      Estonia  2 Potentially Successful
## 3      0 United Kingdom  3 Potentially Successful
## 4      0 United States  0 Potentially Successful
## 5      0 United States  4 Potentially Successful
## 6      0      Chile   2 Potentially Successful
```

```
nrow(df)
```

```
## [1] 32765
```

Country level analysis

Data Manipulation

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.3.3
```

```
## Warning: package 'ggplot2' was built under R version 4.3.3
```

```
## Warning: package 'tibble' was built under R version 4.3.3
```

```
## Warning: package 'tidyr' was built under R version 4.3.3
```

```
## Warning: package 'readr' was built under R version 4.3.3
```

```
## Warning: package 'purrr' was built under R version 4.3.3
```

```
## Warning: package 'dplyr' was built under R version 4.3.3
```

```
## Warning: package 'forcats' was built under R version 4.3.3
```

```
## Warning: package 'lubridate' was built under R version 4.3.3
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats    1.0.0      v stringr    1.5.1
```

```
## v ggplot2    3.5.0      v tibble     3.2.1
```

```
## v lubridate  1.9.3      v tidyr      1.3.1
```

```
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```

top_5_countries <- df %>%
  count(country, sort = TRUE) %>%
  head(5) %>%
  pull(country)

df$country[!(df$country %in% top_5_countries)] <- "Others"
head(df)

```

```

##           name      market funding_total_usd  status
## 1      #waywire      News      1750000  acquired
## 2 'Rock' Your Paper  Publishing      40000  operating
## 3 (In)Touch Network Electronics  1500000  operating
## 4 -R- Ranch and Mine   Tourism      60000  operating
## 5   004 Technologies   Software         0  operating
## 6    1,2,3 Listo  E-Commerce      40000  operating
##           region      city funding_rounds founded_at founded_month
## 1    New York City  New York         1 2012-06-01      2012-06
## 2      Tallinn     Tallinn         1 2012-10-26      2012-10
## 3      London     London          1 2011-04-01      2011-04
## 4    Dallas Fort Worth         2 2014-01-01      2014-01
## 5 Springfield, Illinois Champaign     1 2010-01-01      2010-01
## 6      Santiago Las Condes         1 2012-01-01      2012-01
## founded_quarter founded_year first_funding_at last_funding_at  seed venture
## 1      2012-Q2      2012      2012-06-30      2012-06-30 1750000      0
## 2      2012-Q4      2012      2012-08-09      2012-08-09  40000      0
## 3      2011-Q2      2011      2011-04-01      2011-04-01 1500000      0
## 4      2014-Q1      2014      2014-08-17      2014-09-26      0      0
## 5      2010-Q1      2010      2014-07-24      2014-07-24      0      0
## 6      2012-Q1      2012      2013-02-18      2013-02-18  40000      0
## equity_crowdfunding undisclosed convertible_note debt_financing angel grant
## 1           0           0           0           0           0           0
## 2           0           0           0           0           0           0
## 3           0           0           0           0           0           0
## 4      60000           0           0           0           0           0
## 5           0           0           0           0           0           0
## 6           0           0           0           0           0           0
## private_equity post_ipo_equity post_ipo_debt secondary_market
## 1           0           0           0           0
## 2           0           0           0           0
## 3           0           0           0           0
## 4           0           0           0           0
## 5           0           0           0           0
## 6           0           0           0           0
## product_crowdfunding round_A round_B round_C round_D round_E round_F round_G
## 1           0           0           0           0           0           0           0
## 2           0           0           0           0           0           0           0
## 3           0           0           0           0           0           0           0
## 4           0           0           0           0           0           0           0
## 5           0           0           0           0           0           0           0
## 6           0           0           0           0           0           0           0
## round_H      country age success_metric_updated
## 1      0 United States  2      Successful
## 2      0      Others  2 Potentially Successful

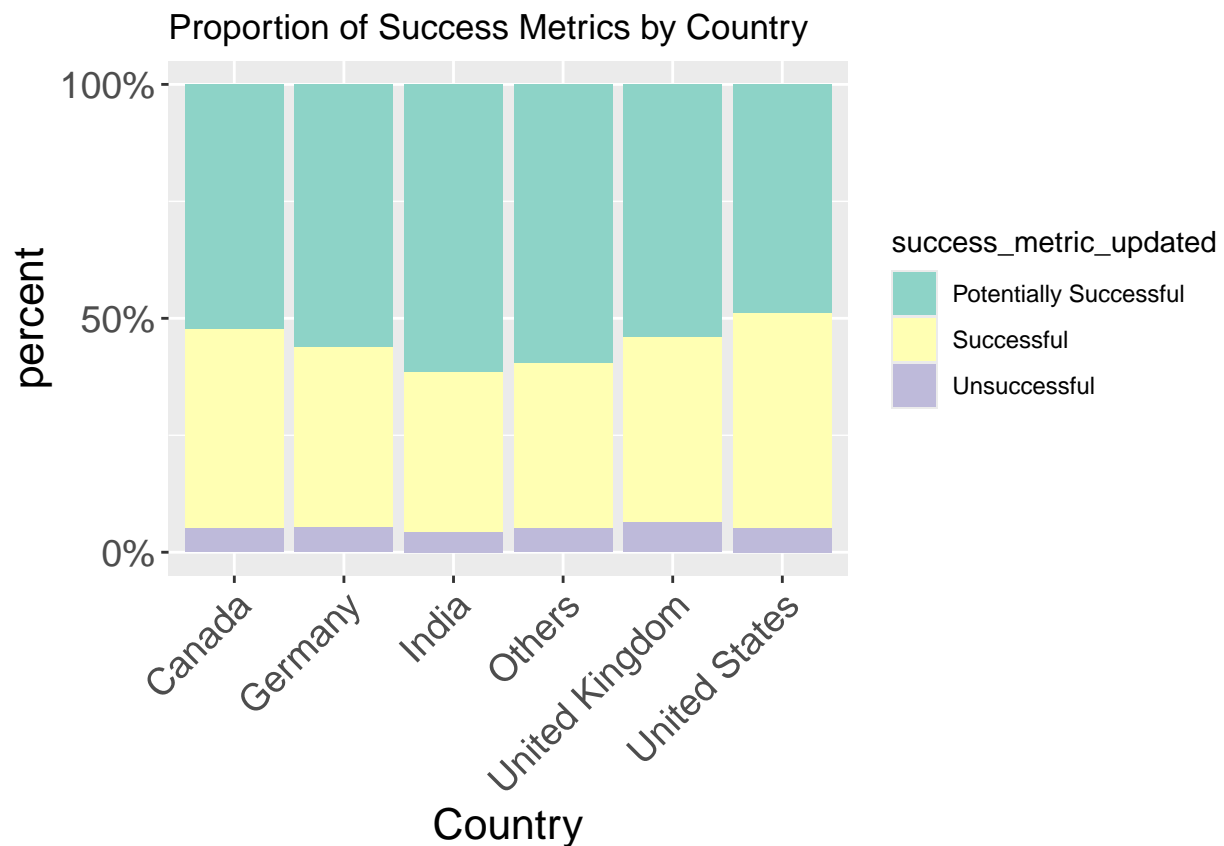
```

```
## 3      0 United Kingdom    3 Potentially Successful
## 4      0 United States    0 Potentially Successful
## 5      0 United States    4 Potentially Successful
## 6      0 Others          2 Potentially Successful
```

Visualising success status proportion (as percentage) for countries

```
library(ggplot2)

ggplot(df, aes(x = country, fill = success_metric_updated)) +
  geom_bar(stat = "count", position = "fill") +
  scale_fill_brewer(palette = "Set3") +
  scale_y_continuous(name = "percent",
                     breaks = c(0, 0.5, 1),
                     labels = scales::percent(c(0, 0.5, 1))) +
  labs(title = "Proportion of Success Metrics by Country",
       x = "Country",
       y = "Proportion") +
  # Increase font size and rotate x-axis labels
  theme(axis.text = element_text(size = 14),
        axis.title = element_text(size = 16),
        axis.text.x = element_text(angle = 45, hjust = 1))
```



Merging “Potentially Successfully” into “Unsuccessful”

```
df_new <- df
df_new$success_metric_updated[df$success_metric_updated == 'Potentially Successful'] <- 'Unsuccessful'
head(df)
```

```
##           name      market funding_total_usd  status
## 1      #waywire      News      1750000  acquired
## 2 'Rock' Your Paper  Publishing      40000  operating
## 3 (In)Touch Network Electronics 1500000  operating
## 4 -R- Ranch and Mine   Tourism      60000  operating
## 5 004 Technologies    Software        0  operating
## 6 1,2,3 Listo E-Commerce      40000  operating
##           region      city funding_rounds founded_at founded_month
## 1 New York City New York      1 2012-06-01      2012-06
## 2 Tallinn Tallinn      1 2012-10-26      2012-10
## 3 London London      1 2011-04-01      2011-04
## 4 Dallas Fort Worth      2 2014-01-01      2014-01
## 5 Springfield, Illinois Champaign      1 2010-01-01      2010-01
## 6 Santiago Las Condes      1 2012-01-01      2012-01
## founded_quarter founded_year first_funding_at last_funding_at seed venture
## 1 2012-Q2 2012 2012-06-30 2012-06-30 1750000 0
## 2 2012-Q4 2012 2012-08-09 2012-08-09 40000 0
## 3 2011-Q2 2011 2011-04-01 2011-04-01 1500000 0
## 4 2014-Q1 2014 2014-08-17 2014-09-26 0 0
## 5 2010-Q1 2010 2014-07-24 2014-07-24 0 0
## 6 2012-Q1 2012 2013-02-18 2013-02-18 40000 0
## equity_crowdfunding undisclosed convertible_note debt_financing angel grant
## 1 0 0 0 0 0 0
## 2 0 0 0 0 0 0
## 3 0 0 0 0 0 0
## 4 60000 0 0 0 0 0
## 5 0 0 0 0 0 0
## 6 0 0 0 0 0 0
## private_equity post_ipo_equity post_ipo_debt secondary_market
## 1 0 0 0 0
## 2 0 0 0 0
## 3 0 0 0 0
## 4 0 0 0 0
## 5 0 0 0 0
## 6 0 0 0 0
## product_crowdfunding round_A round_B round_C round_D round_E round_F round_G
## 1 0 0 0 0 0 0 0
## 2 0 0 0 0 0 0 0
## 3 0 0 0 0 0 0 0
## 4 0 0 0 0 0 0 0
## 5 0 0 0 0 0 0 0
## 6 0 0 0 0 0 0 0
## round_H country age success_metric_updated
## 1 0 United States 2 Successful
## 2 0 Others 2 Potentially Successful
## 3 0 United Kingdom 3 Potentially Successful
## 4 0 United States 0 Potentially Successful
```

```
## 5      0 United States  4 Potentially Successful
## 6      0      Others   2 Potentially Successful
```

```
df_new <- df_new[c("country", "success_metric_updated")]
head(df_new)
```

```
##      country success_metric_updated
## 1 United States      Successful
## 2      Others      Unsuccessful
## 3 United Kingdom      Unsuccessful
## 4 United States      Unsuccessful
## 5 United States      Unsuccessful
## 6      Others      Unsuccessful
```

Contingency tables

```
tbl <- table(df_new$country, df_new$success_metric_updated)
tbl
```

```
##
##      Successful Unsuccessful
## Canada      441      594
## Germany     269      428
## India       220      424
## Others     2266     4145
## United Kingdom 712     1084
## United States 10154    12028
```

```
library(lsr)
```

```
## Warning: package 'lsr' was built under R version 4.3.3
```

```
cramer_v <- crammersV(tbl)
cat("Cramer's V:", cramer_v)
```

```
## Cramer's V: 0.08863159
```

```
tbl_prop <- prop.table(tbl, 1)
tbl_prop
```

```
##
##      Successful Unsuccessful
## Canada      0.4260870  0.5739130
## Germany     0.3859397  0.6140603
## India       0.3416149  0.6583851
## Others     0.3534550  0.6465450
## United Kingdom 0.3964365  0.6035635
## United States 0.4577585  0.5422415
```

Chi square

```
chisq.test(tbl)
```

```
##  
## Pearson's Chi-squared test  
##  
## data:  tbl  
## X-squared = 257.39, df = 5, p-value < 2.2e-16
```

Testing if proportion of success for US is different from other countries

```
# Create a new data frame with the "Others" category  
df_new_merge <- df[c("country", "success_metric_updated")]  
df_new_merge$country[!(df_new_merge$country %in% c("United States", "Others"))] <- "Others"  
  
# Create a new table with the updated data frame  
tbl_merge <- table(df_new_merge$country, df_new_merge$success_metric_updated)  
  
tbl_merge
```

```
##  
##           Potentially Successful Successful Unsuccessful  
## Others           6122           3908           553  
## United States    10879          10154          1149
```

```
cramer_v <- cramerV(tbl_merge)  
cat("Cramer's V:", cramer_v)
```

```
## Cramer's V: 0.08518362
```

```
# Subset data for US and Others  
usa_data <- df_new_merge[df_new_merge$country == "United States", ]  
other_data <- df_new_merge[df_new_merge$country != "United States", ]  
  
# Calculate success proportions for each group  
p_usa <- sum(usa_data$success_metric_updated == "Successful") / nrow(usa_data)  
p_other <- sum(other_data$success_metric_updated == "Successful") / nrow(other_data)  
  
# Calculate pooled proportion (assuming equal variances)  
pooled_p <- (sum(usa_data$success_metric_updated == "Successful") + sum(other_data$success_metric_updated == "Successful")) / (nrow(usa_data) + nrow(other_data))  
  
# Calculate standard error for the difference in proportions  
se_diff <- sqrt(pooled_p * (1 - pooled_p) * (1/nrow(usa_data) + 1/nrow(other_data)))  
  
# Calculate the z-statistic  
z_stat <- (p_usa - p_other) / se_diff
```

```

# Interpretation
# A p-value less than the chosen significance level 0.05 indicates
# a statistically significant difference between the success proportions
# in US and Others category.

# Confidence level (e.g., 95%)
confidence_level <- 0.95

# Critical value (one-tailed) from standard normal distribution table for confidence level
z_crit <- qnorm(1 - (1 - confidence_level) / 2, mean = 0, sd = 1)

# Confidence interval for the difference in proportions
lower_bound <- (p_usa - p_other) - z_crit * se_diff
upper_bound <- (p_usa - p_other) + z_crit * se_diff

# Interpretation
# Report the z-statistic, p-value, and confidence interval.
# The confidence interval captures the range of plausible values for the true
# difference in success proportions between US and Others with a confidence level
# of (e.g.,) 95%.

cat("z-statistic:", z_stat, "\n")

## z-statistic: 15.13249

cat("p-value (two-tailed):", 2 * pnorm(-abs(z_stat), mean = 0, sd = 1, lower.tail = TRUE), "\n")

## p-value (two-tailed): 9.887838e-52

cat("Confidence Interval (", confidence_level * 100, "%):", lower_bound, " - ", upper_bound, "\n")

## Confidence Interval ( 95 %): 0.0770262 - 0.09994794

```

City level analysis

```

df_us <- df %>%
  filter(country == "United States")

length(unique(df_us$city))

## [1] 1975

top_5_cities <- df_us %>%
  count(city, sort = TRUE) %>%
  head(5) %>%
  pull(city)

df_us$city[!(df_us$city %in% top_5_cities)] <- "Others"
head(df_us)

```


##	name	market	funding_total_usd	status				
## 1	#waywire	News	1750000	acquired				
## 2	-R- Ranch and Mine	Tourism	60000	operating				
## 3	004 Technologies	Software	0	operating				
## 4	1-800-DENTIST	Health and Wellness	0	operating				
## 5	1-800-DOCTORS	Health and Wellness	1750000	operating				
## 6	1.618 Technology	Real Estate	0	operating				
##	region	city	funding_rounds	founded_at	founded_month			
## 1	New York City	New York	1	2012-06-01	2012-06			
## 2	Dallas	Others	2	2014-01-01	2014-01			
## 3	Springfield, Illinois	Others	1	2010-01-01	2010-01			
## 4	Los Angeles	Others	1	1986-01-01	1986-01			
## 5	Newark	Others	1	1984-01-01	1984-01			
## 6	Orlando	Others	1	2013-12-07	2013-12			
##	founded_quarter	founded_year	first_funding_at	last_funding_at	seed	venture		
## 1	2012-Q2	2012	2012-06-30	2012-06-30	1750000	0		
## 2	2014-Q1	2014	2014-08-17	2014-09-26	0	0		
## 3	2010-Q1	2010	2014-07-24	2014-07-24	0	0		
## 4	1986-Q1	1986	2010-08-19	2010-08-19	0	0		
## 5	1984-Q1	1984	2011-03-02	2011-03-02	0	0		
## 6	2013-Q4	2013	2014-01-22	2014-01-22	0	0		
##	equity_crowdfunding	undisclosed	convertible_note	debt_financing	angel	grant		
## 1	0	0	0	0	0	0		
## 2	60000	0	0	0	0	0		
## 3	0	0	0	0	0	0		
## 4	0	0	0	0	0	0		
## 5	0	0	1750000	0	0	0		
## 6	0	0	0	0	0	0		
##	private_equity	post_ipo_equity	post_ipo_debt	secondary_market				
## 1	0	0	0	0				
## 2	0	0	0	0				
## 3	0	0	0	0				
## 4	0	0	0	0				
## 5	0	0	0	0				
## 6	0	0	0	0				
##	product_crowdfunding	round_A	round_B	round_C	round_D	round_E	round_F	round_G
## 1	0	0	0	0	0	0	0	0
## 2	0	0	0	0	0	0	0	0
## 3	0	0	0	0	0	0	0	0
## 4	0	0	0	0	0	0	0	0
## 5	0	0	0	0	0	0	0	0
## 6	0	0	0	0	0	0	0	0
##	round_H	country	age	success_metric_updated				
## 1	0	United States	2	Successful				
## 2	0	United States	0	Potentially Successful				
## 3	0	United States	4	Potentially Successful				
## 4	0	United States	28	Successful				
## 5	0	United States	30	Successful				
## 6	0	United States	1	Potentially Successful				

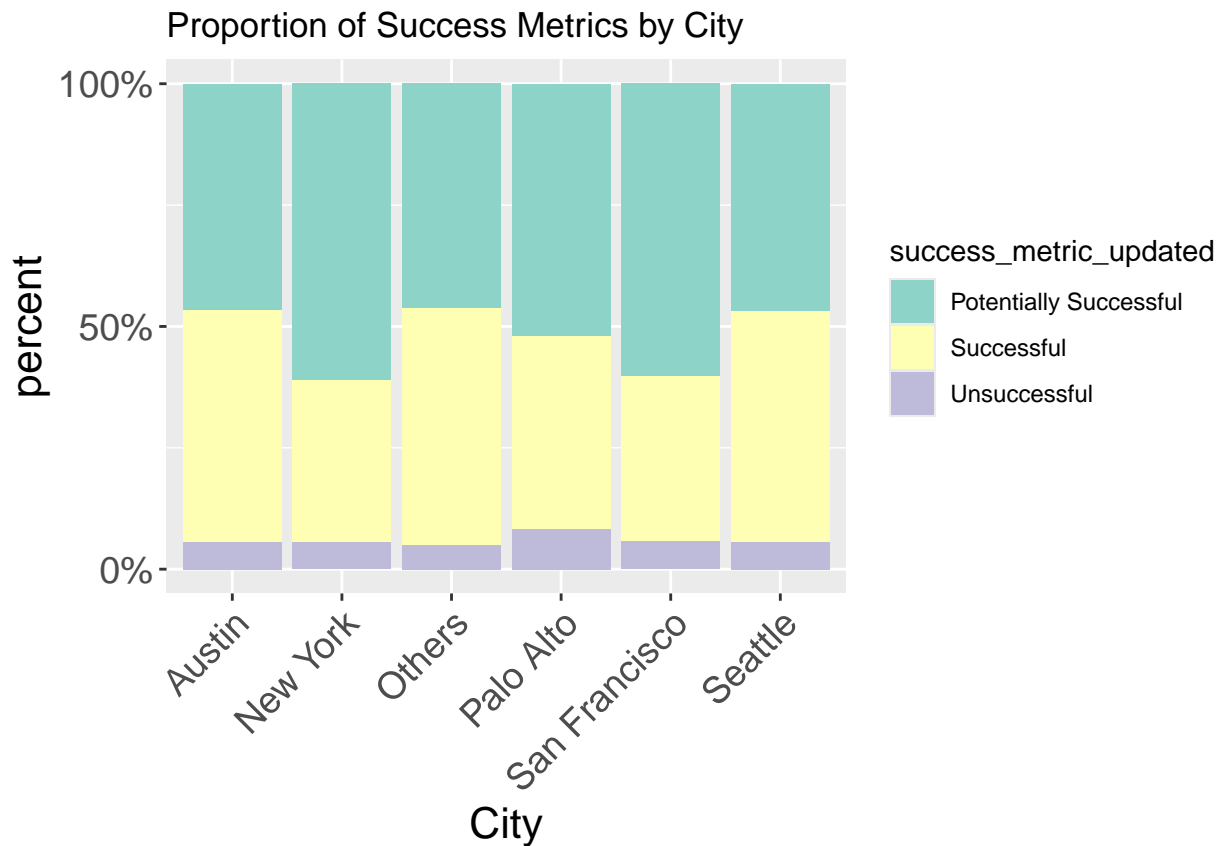
```
library(ggplot2)
```

```
ggplot(df_us, aes(x = city, fill = success_metric_updated)) +  
  geom_bar(stat = "count", position = "fill") +
```

```

scale_fill_brewer(palette = "Set3") +
scale_y_continuous(name = "percent",
                   breaks = c(0, 0.5, 1),
                   labels = scales::percent(c(0, 0.5, 1))) +
labs(title = "Proportion of Success Metrics by City",
     x = "City",
     y = "Proportion") +
# Increase font size and rotate x-axis labels
theme(axis.text = element_text(size = 14), # Change 14 to your desired size
      axis.title = element_text(size = 16), # Change 16 to your desired size
      axis.text.x = element_text(angle = 45, hjust = 1)) # Rotate x-axis labels

```



```

df_us_new <- df_us
df_us_new$success_metric_updated[df_us$success_metric_updated == 'Potentially Successful'] <- 'Unsuccessful'

df_us_new <- df_us_new[c("city", "success_metric_updated")]
head(df_us_new)

```

```

##      city success_metric_updated
## 1 New York      Successful
## 2  Others      Unsuccessful
## 3  Others      Unsuccessful
## 4  Others      Successful
## 5  Others      Successful
## 6  Others      Unsuccessful

```

```
tbl2 <- table(df_us_new$city,df_us_new$success_metric_updated)
tbl2
```

```
##
##           Successful Unsuccessful
## Austin           220           242
## New York         655          1296
## Others           8118          8500
## Palo Alto        192           292
## San Francisco    756          1463
## Seattle          213           235
```

```
cramer_v <- cramersV(tbl2)

cat("Cramer's V:", cramer_v)
```

```
## Cramer's V: 0.1184705
```

```
tbl_prop2 <- prop.table(tbl2, 1)
tbl_prop2
```

```
##
##           Successful Unsuccessful
## Austin      0.4761905    0.5238095
## New York     0.3357253    0.6642747
## Others       0.4885064    0.5114936
## Palo Alto    0.3966942    0.6033058
## San Francisco 0.3406940    0.6593060
## Seattle      0.4754464    0.5245536
```

```
chisq.test(tbl2)
```

```
##
## Pearson's Chi-squared test
##
## data:  tbl2
## X-squared = 311.33, df = 5, p-value < 2.2e-16
```

```
# Create a new data frame with the "Others" category
df_us_new_merge <- df[c("city", "success_metric_updated")]
df_us_new_merge$city[!(df_us_new_merge$city %in% c("Seattle", "Others"))] <- "Others"

# Create a new table with the updated data frame
tbl_us_merge <- table(df_us_new_merge$city, df_us_new_merge$success_metric_updated)

tbl_us_merge
```

```
##
##           Potentially Successful Successful Unsuccessful
## Others           16791          13849          1677
## Seattle           210           213           25
```

```

# Subset data for US and Others
seattle_data <- df_us_new_merge[df_us_new_merge$city == "Seattle", ]
other_data <- df_us_new_merge[df_us_new_merge$city != "Seattle", ]

# Calculate success proportions for each group
p_sea <- sum(seattle_data$success_metric_updated == "Successful") / nrow(seattle_data)
p_other <- sum(other_data$success_metric_updated == "Successful") / nrow(other_data)

# Calculate pooled proportion (assuming equal variances)
pooled_p <- (sum(seattle_data$success_metric_updated == "Successful") + sum(other_data$success_metric_u

# Calculate standard error for the difference in proportions
se_diff <- sqrt(pooled_p * (1 - pooled_p) * (1/nrow(seattle_data) + 1/nrow(other_data)))

# Calculate the z-statistic
z_stat <- (p_sea - p_other) / se_diff

# Interpretation
# A p-value less than the chosen significance level 0.05 indicates
# a statistically significant difference between the success proportions
# in US and Others category.

# Confidence level (e.g., 95%)
confidence_level <- 0.95

# Critical value (one-tailed) from standard normal distribution table for confidence level
z_crit <- qnorm(1 - (1 - confidence_level) / 2, mean = 0, sd = 1)

# Confidence interval for the difference in proportions
lower_bound <- (p_sea - p_other) - z_crit * se_diff
upper_bound <- (p_sea - p_other) + z_crit * se_diff

# Interpretation
# Report the z-statistic, p-value, and confidence interval.
# The confidence interval captures the range of plausible values for the true
# difference in success proportions between US and Others with a confidence level
# of (e.g.,) 95%.

cat("z-statistic:", z_stat, "\n")

## z-statistic: 1.992275

cat("p-value (two-tailed):", 2 * pnorm(-abs(z_stat), mean = 0, sd = 1, lower.tail = TRUE), "\n")

## p-value (two-tailed): 0.04634092

cat("Confidence Interval (", confidence_level * 100, "%):", lower_bound, " - ", upper_bound, "\n")

## Confidence Interval ( 95 %): 0.0007607936 - 0.09305993

```

Additional analysis: Logistic Regression

Country level analysis

```
df_new$success_metric_updated <- as.numeric(df_new$success_metric_updated == "Successful")
```

```
df_new <- within(df_new, relevel(factor(country), ref = "Others"))
head(df_new)
```

```
##           country success_metric_updated
## 1 United States                1
## 2 Others                    0
## 3 United Kingdom              0
## 4 United States              0
## 5 United States              0
## 6 Others                    0
```

```
model <- glm(success_metric_updated ~ country, data = df_new, family = "binomial")
summary(model)
```

```
##
## Call:
## glm(formula = success_metric_updated ~ country, family = "binomial",
##      data = df_new)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.29783    0.06286  -4.738 2.16e-06 ***
## countryGermany -0.16658    0.10002  -1.665 0.095841 .
## countryIndia   -0.35827    0.10419  -3.439 0.000584 ***
## countryOthers  -0.30605    0.06807  -4.496 6.92e-06 ***
## countryUnited Kingdom -0.12250    0.07923  -1.546 0.122091
## countryUnited States  0.12846    0.06429   1.998 0.045681 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 44762  on 32764  degrees of freedom
## Residual deviance: 44502  on 32759  degrees of freedom
## AIC: 44514
##
## Number of Fisher Scoring iterations: 4
```

```
library("effects")
```

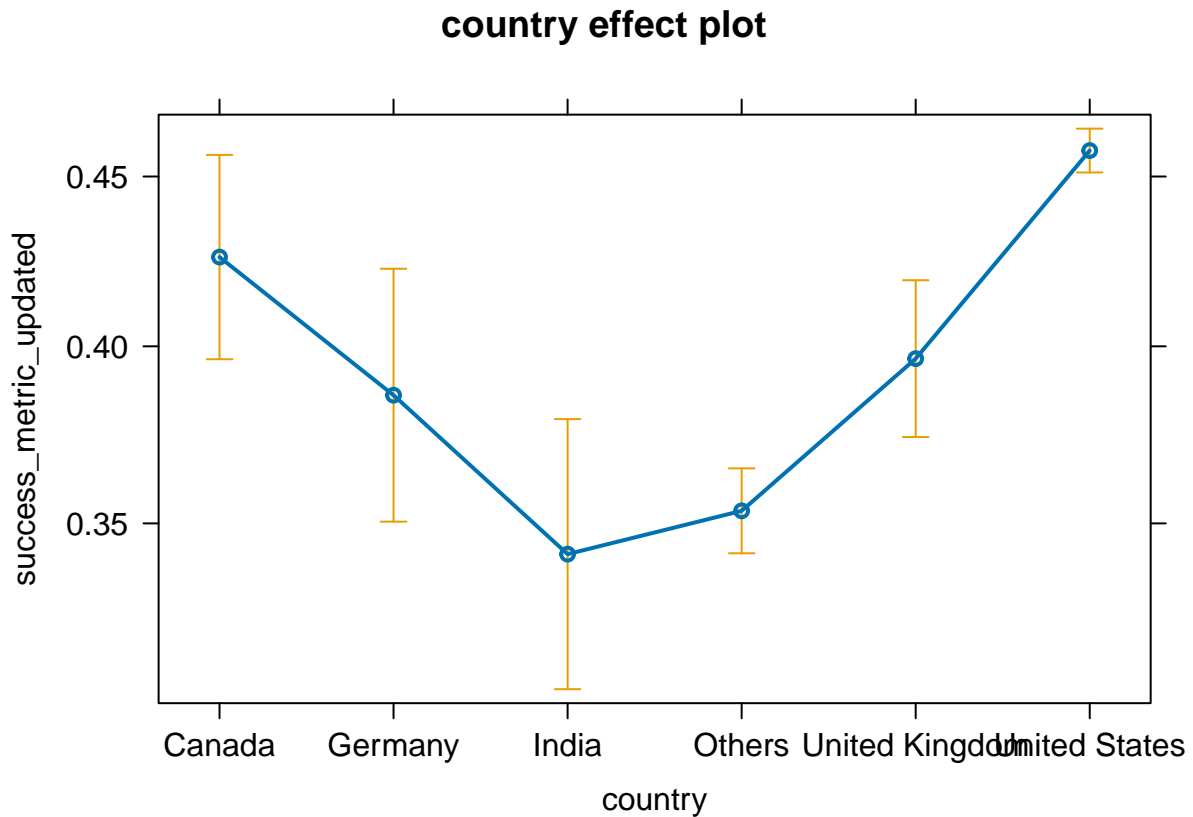
```
## Warning: package 'effects' was built under R version 4.3.3
```

```
## Loading required package: carData
```

```
## Warning: package 'carData' was built under R version 4.3.3
```

```
## lattice theme set by effectsTheme()  
## See ?effectsTheme for details.
```

```
plot(Effect('country', model), confint=T)
```



```
# Get the estimated coefficients  
coefs <- coef(model, simplify=TRUE)  
  
# Calculate odds ratios for each country compared to the reference category  
odds_ratios <- as.data.frame(exp(coefs))  
odds_ratios
```

```
##               exp(coefs)  
## (Intercept)    0.7424242  
## countryGermany 0.8465573  
## countryIndia    0.6988833  
## countryOthers   0.7363482  
## countryUnited Kingdom 0.8847052  
## countryUnited States 1.1370815
```

City level analysis

```
df_us_new <- within(df_us_new, relevel(factor(city), ref = "Others"))
```

```
library(nnet)
df_us_new$city <- as.factor(df_us_new$city)
df_relevel_us <- df_us_new %>%
  mutate(city = relevel(city, ref = "Seattle"))
model2 <- multinom(success_metric_updated ~ city, data = df_relevel_us)
```

```
## # weights:  7 (6 variable)
## initial  value 15375.390759
## iter   10 value 15137.631892
## iter   10 value 15137.631885
## iter   10 value 15137.631885
## final   value 15137.631885
## converged
```

```
summary(model2)
```

```
## Call:
## multinom(formula = success_metric_updated ~ city, data = df_relevel_us)
##
## Coefficients:
##              Values Std. Err.
## (Intercept)    0.098291453 0.09460525
## cityAustin     -0.002981561 0.13276988
## cityNew York    0.584116158 0.10605885
## cityOthers      -0.052309359 0.09586962
## cityPalo Alto   0.320967783 0.13260146
## citySan Francisco 0.561910200 0.10467297
##
## Residual Deviance: 30275.26
## AIC: 30287.26
```

```
df_us_new$success_metric_updated <- as.numeric(df_us_new$success_metric_updated == "Successful")
```

```
model2 <- glm(success_metric_updated ~ city, data = df_us_new, family = "binomial")
summary(model2)
```

```
##
## Call:
## glm(formula = success_metric_updated ~ city, family = "binomial",
##      data = df_us_new)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.095310   0.093154  -1.023   0.3062
## cityNew York  -0.587092   0.104766  -5.604 2.10e-08 ***
## cityOthers     0.049328   0.094438   0.522   0.6014
```

```
## cityPalo Alto      -0.323948   0.131570  -2.462   0.0138 *
## citySan Francisco -0.564893   0.103363  -5.465  4.63e-08 ***
## citySeattle        -0.002983   0.132770  -0.022   0.9821
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 30592  on 22181  degrees of freedom
## Residual deviance: 30275  on 22176  degrees of freedom
## AIC: 30287
##
## Number of Fisher Scoring iterations: 4
```

```
# Get the estimated coefficients
coefs <- coef(model2, simplify=TRUE)

# Calculate odds ratios for each country compared to the reference category
odds_ratios <- as.data.frame(exp(coefs))
odds_ratios
```

```
##              exp(coefs)
## (Intercept)    0.9090909
## cityNew York    0.5559414
## cityOthers      1.0505647
## cityPalo Alto   0.7232877
## citySan Francisco 0.5684211
## citySeattle     0.9970213
```

```
library("effects")
plot(Effect('city', model2), confint=T)
```